A WEB of Sharing Trust–Responsibility Toward a Trustless Future vol.2

Onboarding & Operational Transition for Runtime AI Governance

Second Edition — Version C

Document Control Information

Document Identifier: APOD-TR-004
Title: A WEB of Sharing Trust–Responsibility Toward a Trustless Future Vol.2
Edition: Second Edition — Version C
Date of Issue: 17 January 2026
Status: Public Review Draft (Standards Exploration)
Author: Anthony Siljeg
Affiliation: Alnotrea Labs / Policy of Developments - Access PoD *Initiative.*

Distribution Statement

This publication is released as a **Public Review Draft (Standards**

Exploration). It may be used for:

- academic research and comparative analysis,
- policy development and regulatory consultation,
- standards exploration and working-draft discussion, and
- evaluation of approaches to AI governance that move beyond static, documentation-based compliance toward **engineered, runtime trust mechanisms**.
- Feedback is invited to support refinement prior to subsequent editions.

This document does **not** constitute regulatory guidance, certification, or legal advice. It is intended to support informed review, discussion, and refinement of governance architectures, assurance models, and scoring methodologies for AI systems.

Executive Summary — Version C

Onboarding & Operational Transition for Runtime AI Governance

Volume 2 — Version C addresses a critical and often under-specified phase in AI governance: **onboarding into runtime operation**. While Version A defines the authoritative framework for evidence-based, runtime AI governance, and Version B demonstrates feasibility at execution depth, neither specifies how organizations should *enter* governance responsibly without asserting trust prematurely.

Version C fills this gap.

Purpose and positioning

Version C is a **transitional, operational guidance document**. It does not introduce new governance principles, assurance semantics, or scoring logic. Instead, it clarifies how organizations, regulators, and standards bodies should interpret and manage the **transition from design-time intent to runtime governance**, particularly under conditions of uncertainty.

Version C is explicitly **subordinate to Version A** and **complementary to Version B**. Where ambiguity arises, Version A remains authoritative.

Core insight: containment precedes alignment

Version C is grounded in a clear distinction:

- **Containment** is an operational requirement: the ability to govern AI systems while they are running through runtime controls, evidence generation, escalation, and revocation.
- **Alignment** is a governance maturity outcome that strengthens over time as containment capability, tooling, standards, and professional expertise evolve.

Version C treats containment as a **prerequisite**, not a consequence, of alignment. Alignment cannot be reliably assumed at onboarding and must not be asserted in the absence of operational containment.

What Version C does

Version C:

- defines onboarding as a **governed entry condition**, not an approval or certification event;
- maps onboarding phases to system maturity (under-defined, informally defined, over-defined);
- specifies minimum organizational, technical, and evidentiary conditions for entry into runtime governance;
- clarifies the role of Simulation AI-Powered Policies (SAPP), runtime policy activation, and early-phase escalation;
- sets expectations for regulators and standards bodies during onboarding; and
- identifies common failure modes and governance anti-patterns.

Throughout, Version C emphasizes **evidence over declaration** and **operational behavior over intent**.

What Version C does not do

Version C does **not**:

- certify systems, grant approval, or imply regulatory endorsement;
- define deployment readiness or fitness for purpose;
- prescribe architectures, APIs, configurations, or implementations;
- modify CSC semantics or scoring logic; or
- elevate any reference system, including Q-Pathformer, to a normative requirement.

Its role is to constrain unsafe inference and misuse, not to operationalize systems directly.

Relationship to Q-Pathformer

Q-Pathformer is treated strictly as an **optional, non-normative reference substrate** that may be used during onboarding to exercise and inspect runtime governance mechanisms under controlled conditions. Governance authority, assurance interpretation, and trust signals remain external and independent of any execution substrate.

Organizations are expected to transition beyond reference implementations as governance capability matures, without redefining governance principles.

Why this matters

Onboarding is the **most fragile phase** of AI system operation. It is where governance failure is most likely to be masked by documentation, optimism, or reputational assumptions.

By formalizing onboarding as a governed transition:

- false confidence is reduced;
- certification-by-association is avoided;
- regulators gain clearer visibility into early governance behavior; and
- alignment is allowed to mature safely under containment.

Outcome

Version C makes the Volume 2 framework **operable in practice** without weakening its standards credibility. It provides a disciplined pathway into runtime governance, after which continuous assurance, revocation semantics, and interoperability remain governed exclusively by Version A.

Trust is not granted at entry.
It is engineered, observed, and maintained—continuously.

Second Edition Note — Version C

(Onboarding & Operational Transition)

This Second Edition — **Version C** builds upon **Volume 2 — Version A**, which defines a **model-agnostic, execution-independent AI governance framework** grounded in **continuous, evidence-derived assurance**.

While **Version B** introduces an optional, architecture-specific extension to illustrate how the Version A framework may be instantiated at execution depth, **Version C addresses a different and previously under-specified lifecycle gap**:
how organizations enter runtime AI governance responsibly, without asserting trust prematurely.

Purpose of Version C

Version C introduces **onboarding and operational transition discipline**. It formalizes onboarding as a **governed entry condition**, characterized by:

- containment before alignment,
- evidence generation before assurance,
- and escalation and revocation readiness under uncertainty.

Version C **does not introduce new governance principles**, **does not modify assurance semantics**, and **does not alter Compliance Star Certificate (CSC) scoring logic** defined in Version A.
It exists to make Version A **operable in practice at the point of system entry**, without weakening its standards credibility or authority.

Relationship to Versions A and B

- **Version A** remains the **authoritative reference** for governance authority, terminology, assurance semantics, CSC logic, and conformance interpretation.
- **Version B** remains an **architecture-specific, illustrative extension**, using the Q-Pathformer system solely as a reference implementation to demonstrate execution-level feasibility.
- **Version C** is **subordinate to Version A and complementary to Version B**, focusing exclusively on onboarding and early operational governance under containment.

Where differences or ambiguities arise between editions, **Version A prevails**.

Canonical Artifacts and Release Structure

The **canonical Version A and Version B artifact packs (v1a and v1b)** are published together within the **v1.1-workspace** GitHub release.
This workspace provides an inspection-ready, audit-traceable environment for reviewing authoritative artifacts without modifying their original releases.

The **v1.1-workspace canonical release** is available at:
https://github.com/Access-PoD/access-pod-artifacts/releases/tag/v1.1-workspace

The **Version C paper and associated v1c artifacts**, including onboarding-specific documentation and informative, non-normative companion materials, are published separately as the **v1c release**.

The **v1c release** is available at:
https://github.com/Access-PoD/access-pod-artifacts/releases/tag/v1c

Status and Use

Version C is released as a **Public Review Draft (Standards Exploration)**.

It should be read as **operational transition guidance**, not as:

- a certification scheme,
- a readiness determination,
- a regulatory endorsement, or
- a finalized standards specification.

Version C is shared to support:

- regulator and standards-body review of onboarding discipline,
- organizational interpretation of governed entry conditions,
- identification of governance failure modes at system entry, and
- informed feedback prior to subsequent consolidation or extension (Version D or later).

Trust is **not granted at onboarding**.
It is **engineered, observed, and maintained—continuously**.

Table of Contents

Executive Summary — Version C……2

Edition Note — Version C (Onboarding & Operational Transition)……5
Relationship to Version A and Version B……5

Canonical Artifacts and Release Structure……6

1. **Purpose and Scope of Version C**
   1.1 Why Onboarding Requires Governance Discipline……9
   1.2 What Version C Enables (and What It Explicitly Does Not)……10
   1.3 Intended Audience: Regulators, Enterprises, and System Operators……11
2. **Containment and Alignment in the Context of Onboarding**
   2.1 Containment as an Operational Requirement……12
   2.2 Alignment as a Governance Maturity Outcome……13
   2.3 Why Containment Precedes Alignment……14
   2.4 Alignment Without Containment as a Failure Mode……15
3. **Governance Maturity and Onboarding Phases**
   3.1 Under-Defined Systems: Exploratory Onboarding……16
   3.2 Informally Defined Systems: Provisional Runtime Governance……17
   3.3 Over-Defined Systems: Continuous Assurance……17
   3.4 Entry Conditions for Each Phase……18
   3.5 Prohibited Shortcuts and Common Failure Patterns……18
4. **Organizational Onboarding Model**
   4.1 Required Roles at Onboarding……20
   4.2 Roles That Must Not Be Overloaded……21
   4.3 Separation of Authority Between Governance and Execution……22
   4.4 Human-in-the-Loop Design at Entry……22
5. **Technical Onboarding (Non-Prescriptive)**
   5.1 Identifying Execution Boundaries……24
   5.2 Declaring Enforcement Points……24
   5.3 Telemetry Readiness Checks……25
   5.4 Evidence Continuity Requirements……25
   5.5 Preconditions for Permitting CSC Signals……26
6. **Access PoD in the Onboarding Lifecycle**
   6.1 Simulation AI-Powered Policies (SAPP) at Onboarding……28
   6.2 Runtime Policy Activation……28
   6.3 Escalation and Revocation During Early Operation……29
   6.4 What Onboarding Does Not Certify or Approve……29
7. **Relationship to Q-Pathformer (Informative)**
   7.1 Q-Pathformer as an Onboarding Substrate……31
   7.2 Why Q-Pathformer Is Optional……31

7.3 Using Reference Systems Without Embedding Authority……32
7.4 Transitioning Beyond Reference Implementations……32
8. **Regulator and Standards Engagement During Onboarding**
8.1 What Regulators Can Observe at Each Phase……34
8.2 Evidence Expectations at Entry Versus Maturity……35
8.3 Avoiding Certification-by-Association……35
8.4 Cross-Border and Multi-Sector Considerations……36
9. Failure Modes and Anti-Patterns
9.1 Treating onboarding as approval……37
9.2 Confusing alignment claims with containment……37
9.3 Telemetry gaps at entry……38
9.4 Governance theatre……38
10. **Transition to Continuous Governance**
10.1 From Onboarding to Steady-State Governance……39
10.2 When Version C Hands Back to Version A……39
10.3 Preparing for Version D: Stress Testing and Ethics……40

Conclusion
Governance as a lived capability……41
Why containment enables alignment……41
Why onboarding is the most fragile phase……41

1. Purpose and Scope of Version C

Volume 2 — Version C defines the **onboarding and operational transition context** for the governance framework specified in Version A and instantiated illustratively in Version B. Its purpose is to clarify how organizations may **enter runtime AI governance in a controlled, evidence-based manner**, without treating onboarding as approval, certification, or deployment endorsement.

Version C does not introduce new governance principles, assurance semantics, or scoring logic. Instead, it provides a **governed transition reference** that preserves the separation between governance authority, system execution, and assurance interpretation established in earlier editions.

This section explains why onboarding itself must be treated as a governed activity, what Version C is intended to enable, and who the document is written for.

1.1 Why onboarding requires governance discipline

Onboarding an AI system into operation is a **high-risk transition point** in the system lifecycle. It is the moment at which theoretical governance intent meets real execution behavior, often under incomplete knowledge of downstream impact, emergent behavior, or integration effects.

In many contemporary AI deployments, onboarding is treated as an administrative milestone—characterized by documentation reviews, sign-offs, or policy acknowledgements. Such approaches implicitly assume that governance maturity exists at the point of entry. In practice, this assumption is frequently invalid.

From the perspective of this framework:

- onboarding precedes alignment maturity;
- operational behavior is not yet fully understood; and
- risk profiles are provisional and subject to change.

Without governance discipline at onboarding, organizations may inadvertently:

- assert trust before evidence exists,
- conflate readiness with compliance,
- or allow systems to operate without effective containment mechanisms.

Version C therefore treats onboarding not as an approval event, but as a **governed entry condition**—one in which containment, evidence generation, escalation, and revocation mechanisms must be demonstrably in place before trust can begin to be earned.

1.2 What Version C enables (and what it explicitly does not)

Version C enables reviewers and operators to understand **how runtime governance is entered responsibly**, while maintaining strict scope discipline.

Specifically, Version C enables:

- a shared understanding of onboarding phases aligned to system maturity;
- clarity on when Simulation AI-Powered Policies (SAPP), the Access Policy-on-Demand (PoD) engine, and Compliance Star Certificate (CSC) mechanisms may be activated;
- explicit articulation of authority boundaries during early operation;
- alignment between organizational roles, governance controls, and evidence expectations; and
- regulator- and standards-facing review of transition discipline.

Version C explicitly does **not**:

- imply trust, readiness, or regulatory determination;
- define deployment readiness or fitness for purpose;
- prescribe technical implementations, APIs, or configurations;
- modify governance principles, CSC semantics, or scoring logic defined in Version A; or
- elevate any reference architecture, including Q-Pathformer, to a normative requirement.

By design, Version C avoids implementation prescription. Its function is to constrain unsafe inference and misuse by defining **what onboarding means in governance terms**, not how systems must be built.

Policy Evolution Boundary (Clarification)

Version C recognizes that AI systems may operate within environments where **legacy policies**, **proposed policy instruments**, and **amended or evolving policy frameworks** coexist. Such conditions are common during regulatory transition, standards development, and early system deployment.

However, Version C **does not define**, prescribe, or formalize how policies are drafted, amended, ratified, or superseded. The establishment, amendment, or withdrawal of policy authority

remains external to this document and subject to applicable legal, regulatory, and standards-setting processes.

Within the scope of Version C:

- onboarding **does not constitute policy approval or endorsement**;
- governance mechanisms operate **under existing policy authority**, even where policy evolution is ongoing; and
- uncertainty arising from policy transition **must be treated as a containment condition**, not as a basis for asserting trust or readiness.

Version C is concerned solely with **how systems are onboarded into runtime governance responsibly under conditions of policy uncertainty**, not with determining the legitimacy, maturity, or precedence of policy instruments themselves.

Where policy ambiguity exists, Version C requires **explicit declaration of the governing policy context** and **constrained operation until evidence supports progression**, without substituting onboarding discipline for policy determination.

1.3 Audience: regulators, enterprises, system operators

Version C is intended for stakeholders responsible for **governance accountability at the point of system entry**, including:

- **Regulators and policymakers**, seeking to understand how runtime governance may be introduced prior to full alignment maturity, and how evidence-based oversight can begin without reliance on static approvals.
- **Enterprises and public-sector organizations**, responsible for deploying or operating AI systems under evolving regulatory and societal expectations, and requiring clarity on governance readiness versus deployment endorsement.
- **System operators and governance leads**, tasked with implementing containment mechanisms, maintaining evidence continuity, and managing escalation and revocation during early operational phases.

Version C is not written as a developer manual or an implementation guide. It is a **review and transition document**, intended to support informed oversight, disciplined adoption, and consistent interpretation of onboarding within the broader Volume 2 governance framework.

2. Containment and Alignment in the Context of Onboarding

This section clarifies the distinct but related roles of **containment** and **alignment** during the onboarding phase of AI system operation. It establishes why containment must be treated as an **operational prerequisite**, while alignment is understood as a **governance maturity outcome** that strengthens over time.

Within the Volume 2 framework, onboarding is not the point at which alignment can be reasonably assumed. It is the point at which containment must be demonstrably established so that alignment may evolve safely, measurably, and under continuous oversight.

2.1 Containment as an operational requirement

Containment, within the context of this framework, refers to the **operational ability to govern an AI system while it is running**. It is concerned with whether governance controls are executable, observable, and enforceable under real operating conditions.

Containment is not a theoretical construct. It is defined by the presence of concrete mechanisms that constrain behavior, generate evidence, and enable proportionate intervention when conditions change.

At onboarding, containment is required because system behavior, interaction patterns, and risk exposure are not yet fully characterized. Governance must therefore be capable of responding to uncertainty rather than presuming stability.

Operational containment consists, at a minimum, of the following elements.

Runtime controls

Runtime controls define **where and how governance is enforced during system execution**. These controls operate at identified enforcement points, such as decision gateways, data access interfaces, delegation boundaries, and lifecycle transition events.

Runtime controls ensure that:

- governance constraints are applied concurrently with execution;
- policy violations can be detected as they occur; and
- control failures are observable rather than inferred after the fact.

Without runtime controls, governance remains external to system behavior and cannot function as an operational safeguard.

Evidence generation

Containment requires continuous **generation of governance-relevant evidence**. This evidence must be derived from observed system behavior, not from declarations, documentation, or assumed intent.

Evidence generation enables:

- attribution of actions to defined governance contexts;
- verification that controls are functioning as designed; and
- assessment of whether trust claims remain justified over time.

At onboarding, evidence generation establishes the baseline from which governance maturity may be evaluated. Where evidence does not exist, containment cannot be said to exist.

Escalation and revocation

Containment must include explicit **escalation pathways and revocation mechanisms**. These mechanisms ensure that governance responses scale with risk and that trust is not treated as irrevocable.

Escalation mechanisms:

- route elevated-risk events to appropriate human oversight;
- support contestability and review; and
- preserve accountability across organizational boundaries.

Revocation mechanisms:

- enable suspension or degradation of trust signals when evidence degrades;
- prevent continued operation under invalid assumptions; and
- function as protective controls rather than punitive actions.

Together, escalation and revocation distinguish containment from passive monitoring.

2.2 Alignment as a governance maturity outcome

Alignment refers to the degree to which an AI system's behavior **consistently reflects human, legal, and societal intent** across contexts and over time. Within this framework, alignment is not treated as a binary property or a precondition for onboarding.

Alignment strengthens as governance capability matures. It improves through:

- repeated exposure to real operational conditions;
- refinement of policies and controls;
- accumulation of high-quality evidence; and
- development of specialized expertise across policy, engineering, risk, and oversight roles.

Because alignment depends on observed behavior and institutional learning, it **cannot be reliably asserted at the point of onboarding**.

Why alignment strengthens over time

As systems move from under-defined to informally defined and eventually over-defined maturity stages, governance mechanisms become more effective at:

- detecting drift and misuse;
- adapting controls to emerging patterns; and
- stabilizing expectations around acceptable behavior.

Alignment, in this sense, is an outcome of sustained containment and governance practice, not a prerequisite that can be certified in advance.

Why alignment cannot be assumed at onboarding

At onboarding, uncertainty is inherent. Claims of alignment at this stage are necessarily provisional and risk creating false confidence if not supported by operational evidence.

Assuming alignment prematurely may:

- mask latent risks;
- discourage scrutiny and adaptation; and
- undermine trust when deviations inevitably occur.

Version C therefore treats alignment as **something to be governed toward**, not something to be declared at entry.

2.3 Why containment precedes alignment

Containment must precede alignment because it provides the **conditions under which alignment can be evaluated meaningfully**.

Without containment:

- behavioral drift may go undetected;
- misuse may be identified only after harm occurs; and
- trust may be asserted without evidentiary basis.

Containment enables:

- early detection of deviation;
- controlled experimentation and learning; and
- incremental improvement of governance effectiveness.

By establishing containment at onboarding, organizations create a governance environment in which alignment can strengthen through practice rather than assumption.

2.4 Alignment without containment as a failure mode

Alignment claims made in the absence of operational containment represent a **systemic governance failure mode**.

Such claims often rely on:

- design-time intent;
- policy declarations; or
- static assessments disconnected from runtime behavior.

In practice, alignment without containment:

- cannot be verified continuously;
- cannot respond proportionately to change; and
- cannot support revocation when assumptions fail.

This framework treats alignment as **observable and revisable**, not assertable. Trust is not conferred by intent, reputation, or stated values. It is earned and maintained through containment mechanisms that make alignment observable, contestable, and revisable over time.

3. Governance Maturity and Onboarding Phases

This section describes how AI systems enter and progress through runtime governance in accordance with the **human-centric system maturity model** introduced in Volume 1 and operationalized in Version A.

The maturity stages—**under-defined**, **informally defined**, and **over-defined**—are not abstract classifications. Within Version C, they function as **governed onboarding phases**, each with distinct risk characteristics, governance expectations, and assurance constraints.

Progression between phases shall be supported by **observable evidence**, not by declaration, intent, or documentation alone.

3.1 Under-defined systems: exploratory onboarding

An under-defined system is one whose behavior, boundaries, and downstream impacts are not yet well understood. At onboarding, such systems often appear promising or innovative but exhibit high uncertainty and unquantified risk.

Exploratory onboarding for under-defined systems shall prioritize **containment and learning**, not confidence or approval.

At this phase:

- system behavior is treated as provisional;
- assumptions are explicitly testable; and
- deployment confidence is intentionally constrained.

Governance at this stage shall emphasize:

- simulation-first evaluation through SAPP;
- limited operational exposure;
- strict enforcement of escalation pathways; and
- prohibition of static assurance claims.

No continuous assurance signal shall be issued for under-defined systems beyond exploratory or internal analysis. Onboarding at this stage establishes **eligibility for observation**, not trust.

3.2 Informally defined systems: provisional runtime governance

An informally defined system is one for which there is growing operational understanding and repeatable usage patterns, but where risk remains context-dependent and subject to change.

Provisional onboarding at this stage introduces **runtime governance**, while explicitly recognizing that governance maturity is incomplete.

At this phase:

- runtime controls are active;
- telemetry is continuously generated;
- human-in-the-loop oversight is mandatory for elevated-risk events; and
- assurance outputs, where present, are explicitly marked as provisional.

Governance mechanisms shall support:

- adaptive control thresholds;
- ongoing simulation for newly observed behaviors;
- evidence-driven refinement of policies; and
- early detection of drift or misuse.

Informally defined systems may participate in CSC-based assurance only under constrained scope and with frequent re-evaluation. Trust at this stage is conditional and subject to rapid revision.

3.3 Over-defined systems: continuous assurance

An over-defined system is one whose operating boundaries, expected behaviors, and risk profile are sufficiently understood to support **continuous assurance**.

At this stage, governance maturity enables:

- stable runtime enforcement;
- sustained evidence continuity; and
- lifecycle-aware assurance interpretation.

For over-defined systems:

- runtime telemetry serves as the primary evidence source;
- SAPP is used selectively to validate material changes or context shifts; and

- CSC outputs may be maintained continuously, subject to degradation and revocation semantics.

Continuous assurance does not imply permanence. Trust remains conditional on ongoing evidence quality and control effectiveness. Over-defined status reflects governance capability, not immunity from change.

3.4 Entry conditions for each phase

Entry into each onboarding phase shall be governed by **explicit conditions**, rather than informal judgment or organizational convenience.

At a minimum:

- under-defined onboarding requires simulation capability and escalation readiness;
- informally defined onboarding requires active runtime controls and telemetry continuity; and
- over-defined onboarding requires sustained evidence quality, control stability, and demonstrated governance responsiveness over time.

Progression between phases shall not be accelerated through documentation, intent statements, or reputational claims. Evidence, not aspiration, governs advancement.

3.5 Prohibited shortcuts and common failure patterns

Version C explicitly identifies the following as governance anti-patterns:

- treating onboarding as approval or certification;
- asserting alignment prior to operational containment;
- issuing assurance signals without sufficient telemetry;
- relying on static audits to justify runtime confidence; and
- bypassing human oversight under claims of automation maturity.

Such shortcuts create false confidence and undermine the integrity of governance outcomes. Where they occur, trust signals shall be constrained, suspended, or revoked until corrective evidence is produced.

Governance maturity cannot be compressed through declaration. It must be built, observed, and validated incrementally.

The governance maturity framing applied in this section builds on the human-centric system progression introduced in Volume 1, which describes how emerging systems transition from exploratory uncertainty toward operational stability; Version C applies that progression solely to onboarding discipline and does not redefine its conceptual foundations.

Normative Summary (Section 3)

- AI systems shall be onboarded according to their maturity stage.
- Containment shall precede continuous assurance.
- Progression between onboarding phases shall be evidence-based.
- Provisional trust shall not be represented as stable assurance.
- Governance shortcuts shall be treated as failure conditions.

4. Organizational Onboarding Model

(People, not just systems)

This section defines the **organizational roles, authority boundaries, and human oversight requirements** necessary to support disciplined onboarding into runtime AI governance. It recognizes that effective containment at entry is not achieved through tooling alone, but through **clear allocation of responsibility across specialized roles**.

Within this framework, onboarding is treated as an organizational capability, not a technical configuration task. Governance maturity depends on whether people, processes, and systems are aligned to support continuous evidence generation, escalation, and accountability.

4.1 Required roles at onboarding

Onboarding into runtime governance requires a **minimum set of distinct roles**, each responsible for a different aspect of containment and assurance. These roles may be held by individuals or small teams, but their responsibilities shall remain clearly separated.

Governance architect

The governance architect is responsible for translating governance intent into **executable governance structures**.

This role:

- defines policy-as-code boundaries and constraints;
- aligns governance controls with system maturity stage;
- ensures consistency with Version A governance principles; and
- maintains separation between governance authority and system execution.

The governance architect does not operate the system and does not interpret assurance outcomes. Their responsibility is structural, not operational.

Runtime risk operator

The runtime risk operator is responsible for **day-to-day oversight of governance signals during operation**, particularly during early onboarding phases.

This role:

- monitors telemetry and escalation indicators;

- responds to elevated-risk events;
- coordinates human-in-the-loop review when thresholds are exceeded; and
- ensures that provisional trust is not treated as stable assurance.

The runtime risk operator operates within defined governance constraints and does not modify policy intent unilaterally.

Evidence & assurance lead

The evidence and assurance lead is responsible for **evidence integrity, continuity, and interpretability**.

This role:

- verifies that governance-relevant telemetry is complete and attributable;
- ensures evidence packaging supports internal review and external oversight;
- manages assurance disclosures to regulators or auditors; and
- confirms that CSC outputs, where applicable, are interpreted within scope and validity context.

The evidence and assurance lead does not execute the system and does not define governance policy.

4.2 Roles that must not be overloaded

A common onboarding failure pattern is the **conflation of governance roles**, particularly in early deployments where staffing is limited.

Within this framework:

- governance architects shall not operate production systems;
- runtime risk operators shall not define governance policy; and
- evidence and assurance leads shall not approve execution changes.

Overloading these roles:

- obscures accountability;
- increases the risk of unchallenged assumptions; and
- weakens containment during the most uncertain phase of operation.

Where resource constraints exist, role separation shall be preserved in function, even if individuals temporarily hold multiple responsibilities under explicit procedural safeguards.

4.3 Separation of authority (governance vs execution)

A foundational requirement of Version C onboarding is the **explicit separation between governance authority and execution capability**.

Accordingly:

- governance roles may authorize, constrain, or revoke execution;
- execution roles may not define or override governance constraints; and
- assurance interpretation remains independent of both governance design and system operation.

This separation ensures that:

- no single role can propose, approve, and execute changes unilaterally; and
- trust outcomes remain evidence-derived rather than self-asserted.

Failure to maintain this separation shall be treated as a containment failure, not a procedural exception.

4.4 Human-in-the-loop design at entry

Human oversight is a **mandatory control** during onboarding. Automation maturity does not eliminate the need for human judgment at entry; it increases the importance of structured intervention pathways.

Human-in-the-loop design at onboarding shall ensure that:

- escalation triggers are explicit and observable;
- intervention authority is clearly assigned;
- decisions are traceable through evidence; and
- overrides are logged and reviewable.

At onboarding, human involvement is not a fallback—it is a **first-class governance mechanism** that supports learning, correction, and accountability as containment mechanisms are exercised under real conditions.

Normative Summary (Section 4)

- Organisational onboarding requires defined governance roles.
- Governance, execution, and assurance responsibilities shall remain distinct.
- Role overloading shall be avoided or explicitly mitigated.
- Human-in-the-loop oversight is mandatory at entry.
- Accountability shall be preserved through clear authority boundaries.

5. Technical Onboarding (Non-Prescriptive)

This section defines the **technical readiness conditions** required to support governed onboarding into runtime AI governance. It does not prescribe system architectures, interfaces, or implementation details. Instead, it specifies **what must be observable and controllable** before governance mechanisms may be relied upon.

Technical onboarding, within Version C, is concerned with **governability**, not optimization. A system that cannot be observed, constrained, or evidenced at runtime cannot be governed, regardless of performance or design intent.

5.1 Identifying execution boundaries

Execution boundaries are the points at which system behavior may produce material impact and therefore require governance oversight. Identifying these boundaries is a prerequisite for operational containment.

Execution boundaries may include, but are not limited to:

- decision or inference gateways;
- data access, transformation, or output interfaces;
- delegation to external models, tools, or services;
- lifecycle transitions such as retraining, fine-tuning, or configuration change; and
- human interaction points with potential downstream effect.

At onboarding, organizations shall be able to **identify and enumerate execution boundaries** relevant to the system's current maturity stage. Where boundaries cannot be identified, governance enforcement cannot be meaningfully applied.

5.2 Declaring enforcement points

Enforcement points are execution boundaries at which governance controls may be applied. Declaring enforcement points does not require implementation detail, but it does require clarity of intent and scope.

For onboarding purposes:

- enforcement points shall be explicitly declared;
- each enforcement point shall be attributable to a governance context; and
- enforcement points shall support observable outcomes (e.g., allow, block, escalate).

Enforcement points establish where governance authority may act. Without declared enforcement points, containment remains theoretical and cannot be exercised under real conditions.

5.3 Telemetry readiness checks

Telemetry readiness refers to the system's ability to produce **governance-relevant signals** during operation. Telemetry is required to support evidence-based assurance and lifecycle evaluation.

At onboarding, telemetry readiness shall be assessed by confirming that:

- governance-relevant events are observable;
- telemetry is attributable to defined enforcement points;
- signals are timestamped and contextualized; and
- telemetry generation is continuous rather than episodic.

Telemetry readiness does not require that all metrics be finalized or optimized. It requires only that **sufficient signals exist to support containment, review, and escalation**. Where telemetry is absent or unreliable, assurance claims shall not be made.

5.4 Evidence continuity requirements

Evidence continuity ensures that governance-relevant evidence remains coherent across system lifecycle events, including updates, retraining, delegation, and scaling.

For onboarding:

- evidence shall persist across execution sessions;
- changes to system state shall not invalidate prior evidence without traceability;
- evidence gaps shall be detectable; and
- loss of evidence continuity shall be treated as a governance event.

Evidence continuity is necessary to prevent false confidence arising from partial visibility or selective observation. Where evidence continuity cannot be maintained, trust signals shall be constrained or withheld.

5.5 What must exist before CSC signals are permitted

The Compliance Star Certificate (CSC) is a **continuous assurance signal**, not an onboarding artifact. CSC signals shall not be issued or interpreted until minimum technical conditions are met.

Dataset Status Declaration and Boundary Conditions

During onboarding, AI systems may rely on datasets that exist in different states, including **current operational datasets**, **proposed or experimental datasets under testing**, and **datasets pending formal verification or approval under applicable governance regimes**.

Version C **does not define dataset validation, certification, or approval criteria**, nor does it establish dataset lifecycle stages as a normative construct. Dataset verification authority, where applicable, is governed exclusively under Version A assurance mechanisms or external regulatory and institutional processes.

For the purposes of onboarding under Version C:

- dataset **status must be explicitly declared** at entry;
- dataset inclusion **shall not be interpreted as evidence of dataset validity, quality, or approval**; and
- dataset-derived outputs **shall not be used to assert trust, readiness, or compliance** in the absence of sufficient operational evidence.

Where datasets are proposed, experimental, or otherwise provisional:

- governance controls shall treat associated system behavior as **conditionally constrained**;
- assurance signals, if generated, shall be **explicitly marked as provisional and non-certifying**; and
- progression toward continuous assurance shall remain dependent on **telemetry-derived evidence**, not dataset designation.

Version C therefore treats dataset status as a **declared input to containment**, not as a trust signal. Dataset presence alone shall not elevate governance posture or justify CSC interpretation.

Before CSC signals are permitted:

- execution boundaries and enforcement points shall be declared;
- telemetry readiness shall be demonstrated;
- evidence continuity shall be established; and
- escalation and revocation mechanisms shall be operational.

Where these conditions are not met:

- CSC outputs shall not be generated, or
- any exploratory signals shall be explicitly marked as non-assurance.

This constraint prevents certification-by-association and ensures that CSC remains evidence-derived rather than inferred.

Normative Summary (Section 5)

- Technical onboarding shall define conditions, not implementations.
- Execution boundaries shall be identifiable.
- Enforcement points shall be declared.
- Telemetry readiness shall be demonstrable.
- Evidence continuity shall be maintained.
- CSC signals shall not be permitted without sufficient operational evidence.

6. Access PoD in the Onboarding Lifecycle

This section describes how the **Access Policy-on-Demand (PoD)** framework operates during onboarding to support **containment, evidence generation, and controlled transition** into runtime governance. It clarifies the role of Simulation AI-Powered Policies (SAPP), runtime policy activation, and early-phase escalation, while explicitly constraining what onboarding does not imply.

Access PoD does not confer trust at onboarding. It enables **governed exposure**, allowing organizations to observe behavior, test assumptions, and intervene proportionately as evidence accumulates.

6.1 SAPP at onboarding (simulation-first containment)

Simulation AI-Powered Policies (SAPP) provide the **primary containment mechanism at onboarding**, particularly for under-defined and informally defined systems.

At onboarding, SAPP shall be used to:

- evaluate proposed execution paths under simulated conditions;
- explore misuse, edge cases, and adversarial scenarios;
- test governance assumptions prior to runtime reliance; and
- generate preliminary evidence to inform whether and how operation may proceed.

Simulation-first containment ensures that novelty, apparent usefulness, or deployment pressure do not bypass governance validation. SAPP outputs at this stage are **informative**, not certifying. They support learning, constraint refinement, and human review before confidence is asserted.

Where simulation outcomes indicate elevated or poorly understood risk, onboarding shall remain constrained until further evidence is produced.

6.2 Runtime policy activation

Runtime policy activation marks the transition from simulation-only governance to **concurrent governance during execution**. This activation shall occur only after minimum onboarding conditions defined in Section 5 are met.

At onboarding, runtime policy activation:

- enforces declared governance constraints at defined enforcement points;
- ensures governance decisions operate concurrently with system execution; and
- enables real-time detection of policy deviation or control failure.

Runtime policy activation does not imply system maturity or stability. During early operation, policies may be adaptive, conservative, and subject to frequent refinement. Governance effectiveness at this stage is evaluated by **observed control behavior**, not by absence of incidents.

6.3 Escalation and revocation during early operation

Early operation is characterized by uncertainty. Accordingly, **escalation and revocation mechanisms shall be treated as first-class controls during onboarding**.

During this phase:

- escalation thresholds shall be conservative and explicit;
- human-in-the-loop intervention shall be readily available;
- governance decisions shall be traceable and reviewable; and
- revocation or degradation of trust signals shall be expected governance outcomes, not exceptions.

Revocation during onboarding is not punitive. It is a protective response indicating that assumptions have been invalidated or that evidence is insufficient to sustain prior trust claims.

Access PoD ensures that escalation and revocation are **operational actions**, not retrospective judgments, allowing containment to remain effective as conditions evolve.

6.4 What onboarding does not certify or approve

Version C explicitly constrains the interpretation of onboarding activities. Onboarding under Access PoD does **not**:

- certify system safety, alignment, or compliance;
- approve deployment, scale-up, or fitness for purpose;
- grant regulatory endorsement or legal standing;
- assert alignment maturity; or
- imply permanence of trust.

Onboarding establishes **eligibility for observation and governance**, not entitlement to trust. Any assurance signals produced during onboarding shall be provisional, scope-bound, and subject to degradation or withdrawal as evidence evolves.

This constraint is essential to prevent certification-by-association and to preserve the integrity of evidence-based governance.

Normative Summary (Section 6)

- SAPP shall be the primary containment mechanism at onboarding.
- Runtime policy activation shall occur only after minimum readiness conditions are met.
- Escalation and revocation shall be expected during early operation.
- Trust signals at onboarding shall be provisional and revisable.
- Onboarding shall not be interpreted as certification, approval, or endorsement.

7. Relationship to Q-Pathformer

(Informative)

This section clarifies the relationship between the Volume 2 governance framework and **Q-Pathformer** during onboarding. It explains how Q-Pathformer may be used as a **reference onboarding substrate**, why its use is **optional**, and how organizations may transition beyond reference implementations without altering governance authority or assurance semantics.

Nothing in this section modifies the normative requirements defined elsewhere in Volume 2.

7.1 Q-Pathformer as an onboarding substrate

Q-Pathformer may be used as an **onboarding substrate** to support early validation of runtime governance concepts under controlled conditions. In this role, Q-Pathformer provides a concrete execution environment in which containment mechanisms—such as enforcement points, telemetry generation, and escalation pathways—can be exercised and observed.

As an onboarding substrate, Q-Pathformer is used to:

- demonstrate how governance controls attach to execution boundaries;
- support simulation-first containment and early runtime governance;
- enable inspection and review of evidence continuity; and
- facilitate standards and regulatory exploration without production-scale deployment.

Use of Q-Pathformer at onboarding is intended to support **learning and validation**, not to establish normative requirements or preferred architectures.

7.2 Why Q-Pathformer is optional

The Volume 2 governance framework is **architecture-agnostic and model-provider-neutral**. Accordingly, use of Q-Pathformer is **not required** for onboarding into runtime governance.

Q-Pathformer is optional because:

- governance authority resides in Version A mechanisms, not in any execution substrate;
- containment and assurance depend on observability and control, not on specific architectures; and
- equivalent execution boundaries may exist in diverse systems with different internal designs.

Organizations may onboard using alternative systems or platforms, provided that the readiness conditions defined in Sections 3–6 are met. No governance outcome, assurance interpretation, or conformance claim depends on the use of Q-Pathformer.

7.3 Using reference systems without embedding authority

Reference systems, including Q-Pathformer, shall not be treated as sources of governance authority, approval, or trust.

When used during onboarding:

- reference systems illustrate feasibility, not compliance;
- governance decisions remain external and independent;
- assurance signals are computed outside the execution substrate; and
- execution environments do not interpret or assert trust.

Embedding governance authority into reference systems would undermine portability, comparability, and interoperability. Version C therefore maintains a strict separation between **governance intent**, **execution substrate**, and **assurance interpretation**, consistent with Versions A and B.

7.4 Transitioning beyond reference implementations

Onboarding using a reference system is not an end state. Organizations are expected to **transition beyond reference implementations** as governance capability matures and operational contexts evolve.

Such transitions:

- do not require redefinition of governance principles;
- do not alter CSC semantics or scoring logic; and
- do not imply loss of conformance or assurance continuity.

Where execution substrates change, governance outcomes shall be reassessed through evidence, not inferred from prior association with a reference system. Trust remains conditional on **current observability, containment, and evidence quality**, regardless of the platform in use.

This approach ensures that governance scales with system evolution while remaining independent of any single technology or implementation.

Informative Summary (Section 7)

- Q-Pathformer may be used as an onboarding substrate.
- Use of Q-Pathformer is optional and non-normative.
- Governance authority shall not be embedded in reference systems.
- Assurance semantics remain external to execution substrates.
- Organizations shall be able to transition beyond reference implementations without redefining governance.

8. Regulator and Standards Engagement During Onboarding

This section describes how regulators and standards bodies may **engage with and observe onboarding activities** under the Volume 2 governance framework. It clarifies what may reasonably be observed at different onboarding phases, how evidence expectations evolve with governance maturity, and how misinterpretation of onboarding as certification is explicitly avoided.

Engagement during onboarding is intended to support **informed oversight and standards exploration**, not to substitute for sector-specific regulatory processes or legal determinations.

8.1 What regulators can observe at each phase

Regulatory and standards engagement during onboarding is necessarily **phase-dependent**, reflecting system maturity and governance readiness.

For **under-defined systems**, regulators may observe:

- simulation-based risk discovery activities;
- declared execution boundaries and enforcement intent;
- existence of escalation and human oversight pathways; and
- evidence that onboarding is constrained and exploratory.

At this phase, regulators should not expect stable assurance outputs or claims of alignment maturity.

For **informally defined systems**, regulators may observe:

- active runtime governance controls at declared enforcement points;
- continuous telemetry generation and attribution;
- provisional assurance signals, where present, clearly marked as such; and
- documented responses to elevated-risk events or control deviations.

For **over-defined systems**, regulators may observe:

- sustained evidence continuity across lifecycle events;
- stable application of runtime controls;
- continuous assurance outputs derived from telemetry; and
- observable degradation or revocation of trust where evidence warrants.

At no phase does onboarding itself constitute regulatory approval or compliance determination.

8.2 Evidence expectations at entry vs maturity

Evidence expectations evolve as systems progress through onboarding phases. Version C distinguishes clearly between **entry evidence** and **maturity evidence**.

At entry:

- evidence demonstrates containment capability;
- telemetry may be incomplete but must be attributable;
- assurance outputs, if any, are provisional; and
- human oversight is central.

At maturity:

- evidence supports continuous assurance;
- telemetry is stable, complete, and lifecycle-aware;
- trust signals are continuously recalculated; and
- governance responses to drift are observable.

This distinction prevents premature assurance claims and supports proportionate oversight aligned to system maturity.

8.3 Avoiding certification-by-association

A recurring risk during onboarding is **certification-by-association**, whereby participation in governance frameworks, pilot programs, or standards exploration is misinterpreted as endorsement or approval.

Version C explicitly rejects certification-by-association. Accordingly:

- onboarding activities shall not be represented as regulatory approval;
- use of reference systems shall not imply compliance;
- provisional assurance outputs shall not be marketed as certification; and
- standards engagement shall not be conflated with conformance claims.

Where misrepresentation occurs, trust signals shall be constrained or withdrawn until corrective clarification is made.

8.4 Cross-border and multi-sector considerations

AI systems frequently operate across jurisdictions and sectors, each with distinct regulatory expectations and risk tolerances. Onboarding must therefore remain **context-aware and adaptable**.

During onboarding:

- sector-specific risk profiles shall be declared where applicable;
- jurisdictional scope of evidence shall be explicit;
- assurance interpretation shall be bounded by declared context; and
- cross-border deployments shall not assume uniform regulatory acceptance.

Standards bodies may use onboarding observations to inform guidance development, but such observations do not override sectoral or jurisdictional requirements.

This approach supports interoperability of governance evidence while preserving regulatory sovereignty.

Normative Summary (Section 8)

- Regulatory observation shall be phase-appropriate.
- Evidence expectations shall scale with governance maturity.
- Onboarding shall not be interpreted as certification or approval.
- Certification-by-association shall be explicitly avoided.
- Cross-border and sectoral context shall be declared and respected.

9. Failure Modes and Anti-Patterns

This section identifies **common failure modes and governance anti-patterns** observed during AI system onboarding. These failures do not arise from malicious intent; they arise from misinterpretation, shortcutting, or overextension of governance concepts beyond what operational evidence can support.

Version C explicitly documents these patterns to prevent misuse of the framework and to reinforce the principle that **trust must be earned through containment and evidence, not inferred from participation or intent**.

9.1 Treating onboarding as approval

A primary failure mode is treating onboarding into a governance framework as an implicit form of approval, certification, or deployment readiness.

Onboarding establishes **eligibility for governed observation**, not endorsement. When onboarding is misrepresented as approval:

- risk is obscured rather than managed;
- scrutiny is reduced at the point of highest uncertainty; and
- accountability becomes diffused.

This framework explicitly rejects onboarding-as-approval. Any representation of onboarding as certification, compliance, or regulatory endorsement constitutes a governance failure.

9.2 Confusing alignment claims with containment

Another common anti-pattern is asserting alignment claims in the absence of operational containment.

Alignment describes the consistency of system behavior with human, legal, and societal intent over time. Containment describes the **ability to govern behavior while the system is running**.

When alignment is claimed without containment:

- behavior cannot be continuously verified;
- drift and misuse may go undetected;
- and confidence becomes narrative-driven rather than evidence-based.

This framework treats containment as a prerequisite for meaningful alignment evaluation. Alignment without containment is not a sign of maturity; it is a source of false confidence.

9.3 Telemetry gaps at entry

Telemetry gaps at onboarding represent a critical failure condition.

Where governance-relevant telemetry is incomplete, delayed, or non-attributable:

- assurance claims cannot be sustained;
- escalation decisions lack evidentiary basis; and
- trust signals, if present, are misleading.

Telemetry gaps shall not be excused as temporary implementation limitations. At onboarding, incomplete telemetry must result in constrained operation and provisional trust, not implicit acceptance.

9.4 Governance theatre

Governance theatre occurs when governance artifacts, committees, or declarations exist without operational effect.

Examples include:

- policies that are not enforceable at runtime;
- assurance statements unsupported by telemetry;
- oversight bodies without escalation authority; and
- documentation substituting for evidence.

Governance theatre creates the appearance of control while masking real risk. Within this framework, governance that does not alter system behavior under real conditions is not governance—it is performative compliance.

Normative Summary (Section 9)

- Onboarding shall not be treated as approval or certification.
- Alignment claims without containment are invalid.
- Telemetry gaps at entry constitute governance failure.
- Governance without operational effect is governance theatre.
- Trust shall not be asserted without evidence.

10. Transition to Continuous Governance

This section describes how organizations move from **onboarding under Version C** to **steady-state governance under Version A**, and how future extensions may be introduced without disrupting governance authority or assurance semantics.

Version C is transitional by design. Its role concludes once containment, evidence continuity, and operational discipline are sufficient to support continuous assurance.

10.1 From onboarding to steady-state

The transition from onboarding to steady-state governance occurs when runtime governance mechanisms operate reliably under normal and stressed conditions.

Indicators that an organization is ready to transition include:

- sustained operation of runtime controls at declared enforcement points;
- continuous, attributable telemetry across lifecycle events;
- effective escalation and intervention under elevated-risk conditions; and
- evidence that trust degradation or revocation has occurred appropriately when assumptions failed.

At this point, governance emphasis shifts from entry discipline to **ongoing assurance maintenance**. Simulation, runtime enforcement, and evidence interpretation continue, but onboarding constraints are relaxed in favor of continuous evaluation.

Steady-state governance does not imply reduced scrutiny. It reflects improved governance capability and confidence grounded in evidence rather than assumption.

10.2 When Version C hands back to Version A

Version C hands back authority to Version A when onboarding considerations no longer dominate governance decisions.

At that point:

- governance intent, assurance semantics, and CSC scoring are governed exclusively by Version A;

- onboarding-specific constraints cease to apply; and
- the system is evaluated primarily through continuous assurance mechanisms.

Version C prepares organizations to operate within Version A by ensuring that **trust remains conditional, evidence-derived, and lifecycle-aware**.

Where uncertainty re-emerges—such as through major system changes, new deployment contexts, or significant drift—organizations may re-enter onboarding-style containment without redefining governance principles.

10.3 Preparing for Version D (stress testing, ethics)

Future governance evolution is expected to focus on **structured stress testing and ethical evaluation** beyond nominal operation.

A subsequent Version D is envisaged to address:

- governance stress testing under adversarial or extreme conditions;
- ethical boundary exploration within controlled parameters; and
- systematic identification of governance failure modes prior to real-world impact.

Version D activities, where introduced, shall remain subordinate to the governance framework defined in Version A and shall remain subordinate to the governance authority and assurance semantics defined in Version A.

Preparation for such future extensions requires only that organizations maintain:

- evidence continuity;
- containment readiness; and
- willingness to re-enter constrained governance modes when conditions change.

Normative Summary (Section 10)

- Transition to steady-state governance shall be evidence-based.
- Version C is transitional and hands authority back to Version A.
- Continuous governance remains conditional and revocable.
- Re-entry into onboarding containment may occur when uncertainty increases.
- Future extensions shall not modify core governance authority or semantics.

Conclusion

Governance as a lived capability

Version C positions governance not as a static requirement or one-time assessment, but as a **lived operational capability**. Effective AI governance is demonstrated through continuous containment, evidence generation, and accountable intervention—not through declarations of intent or compliance artifacts alone.

By formalizing onboarding as a governed transition, Version C enables organizations to enter runtime governance without asserting trust prematurely, while preserving the integrity of the framework defined in Version A.

Why containment enables alignment

Alignment cannot be meaningfully evaluated without containment. Claims of alignment depend on whether system behavior can be constrained, observed, and corrected under real operating conditions.

Version C reinforces that **containment precedes alignment**. It establishes the conditions under which alignment may strengthen over time—through governance maturity, operational learning, and specialization across policy, engineering, and oversight roles. Alignment improves not because systems are assumed to be safe, but because governance is capable of responding when they are not.

Why onboarding is the most fragile phase

Onboarding is the point of greatest uncertainty and the highest risk of false confidence. It is where intent first meets execution, and where governance failure is most likely to be masked by documentation, reputation, or optimism.

Version C treats onboarding as a **fragile but necessary phase**, requiring explicit discipline, constrained trust, and heightened oversight. By doing so, it prevents governance shortcuts and ensures that trust, once earned, remains grounded in evidence rather than assumption.

Closing statement

Version C completes the Volume 2 framework by making **Version A operable in practice**, without weakening its standards credibility or authority. It provides a controlled pathway into runtime governance, after which continuous assurance, revocation semantics, and interoperability remain governed exclusively by Version A.

Trust is not granted at entry. It is engineered, observed, and maintained—continuously.